

CUDA Programming Model Overview

Peter Messmer



3 Ways to Program GPUs



Applications

Libraries

“Drop-in”
Acceleration

OpenACC
Directives

Easily Accelerate
Applications

Programming
Languages

Maximum
Flexibility

Outline



- **CUDA model**
- **CUDA programming basics**
- **Development Resources**
- **GPU architecture for computing**

CUDA Model



What is CUDA?

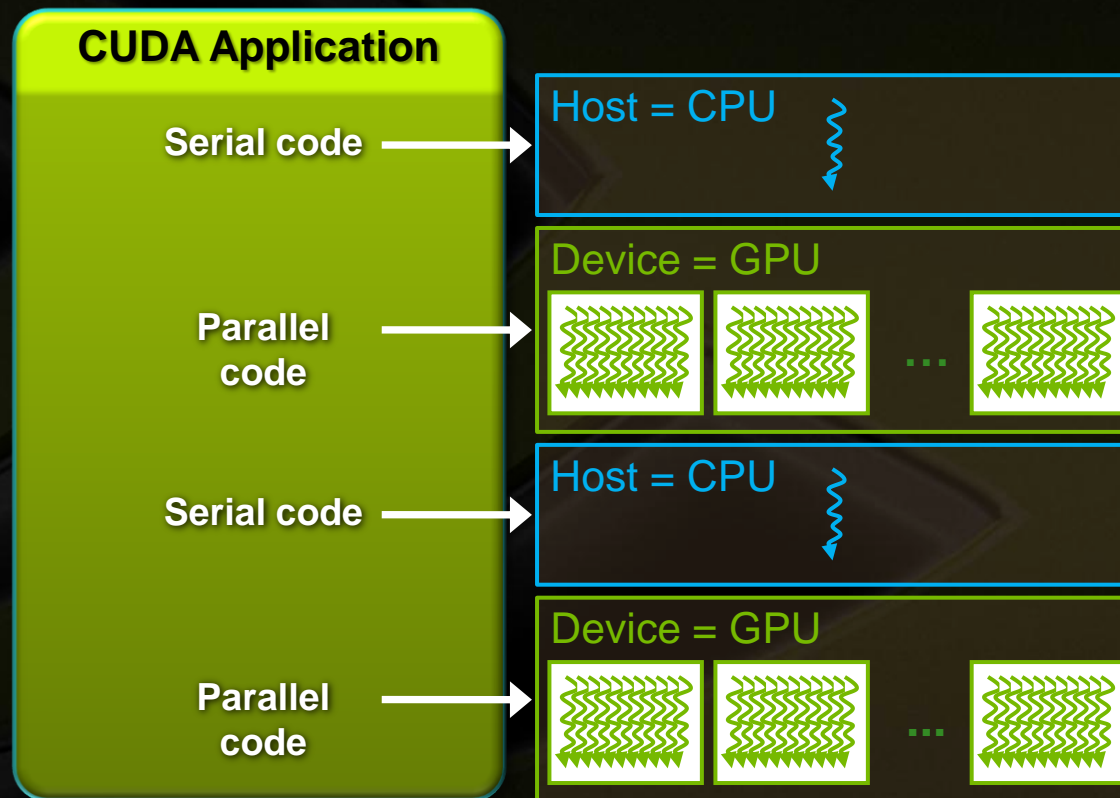


- **C++ with extensions**
 - Fortran support via e.g. PGI's CUDA Fortran
- **CUDA goals:**
 - Scale to 100's of cores, 1000's of parallel threads
 - Let programmers focus on parallel algorithms
 - Enable heterogeneous systems (i.e., CPU+GPU)
- **CUDA defines:**
 - Programming model
 - Memory model

Anatomy of a CUDA Application



- **Serial** code executes in a **Host** (CPU) thread
- **Parallel** code executes in many **Device** (GPU) threads across multiple processing elements



CUDA Kernels



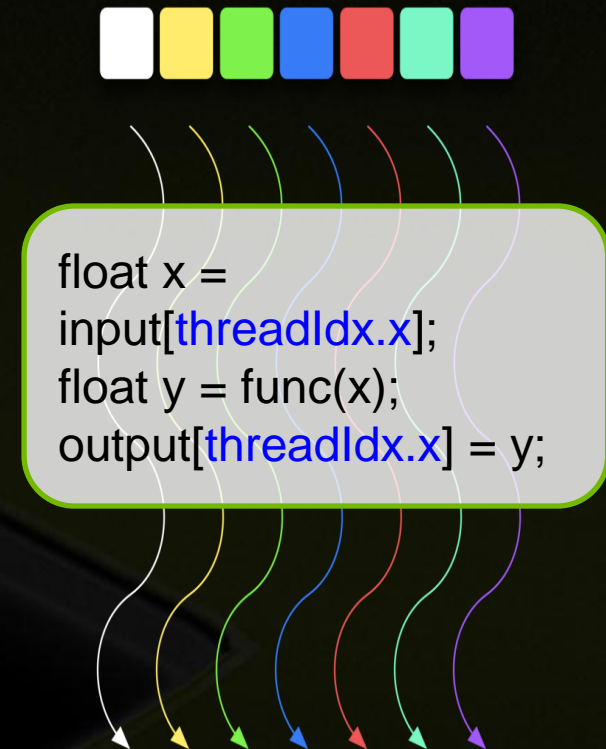
- **Parallel portion of application: execute as a **kernel****
 - Entire GPU executes kernel, many threads
- **CUDA threads:**
 - Lightweight
 - Fast switching
 - 1000s execute simultaneously

CPU	Host	Executes functions
GPU	Device	Executes kernels


CUDA Kernels: Parallel Threads




- A **kernel** is a function executed on the GPU as an array of threads in parallel
- All threads execute the same code, can take different paths
- Each thread has an ID
 - Select input/output data
 - Control decisions



CUDA Kernels: Subdivide into Blocks

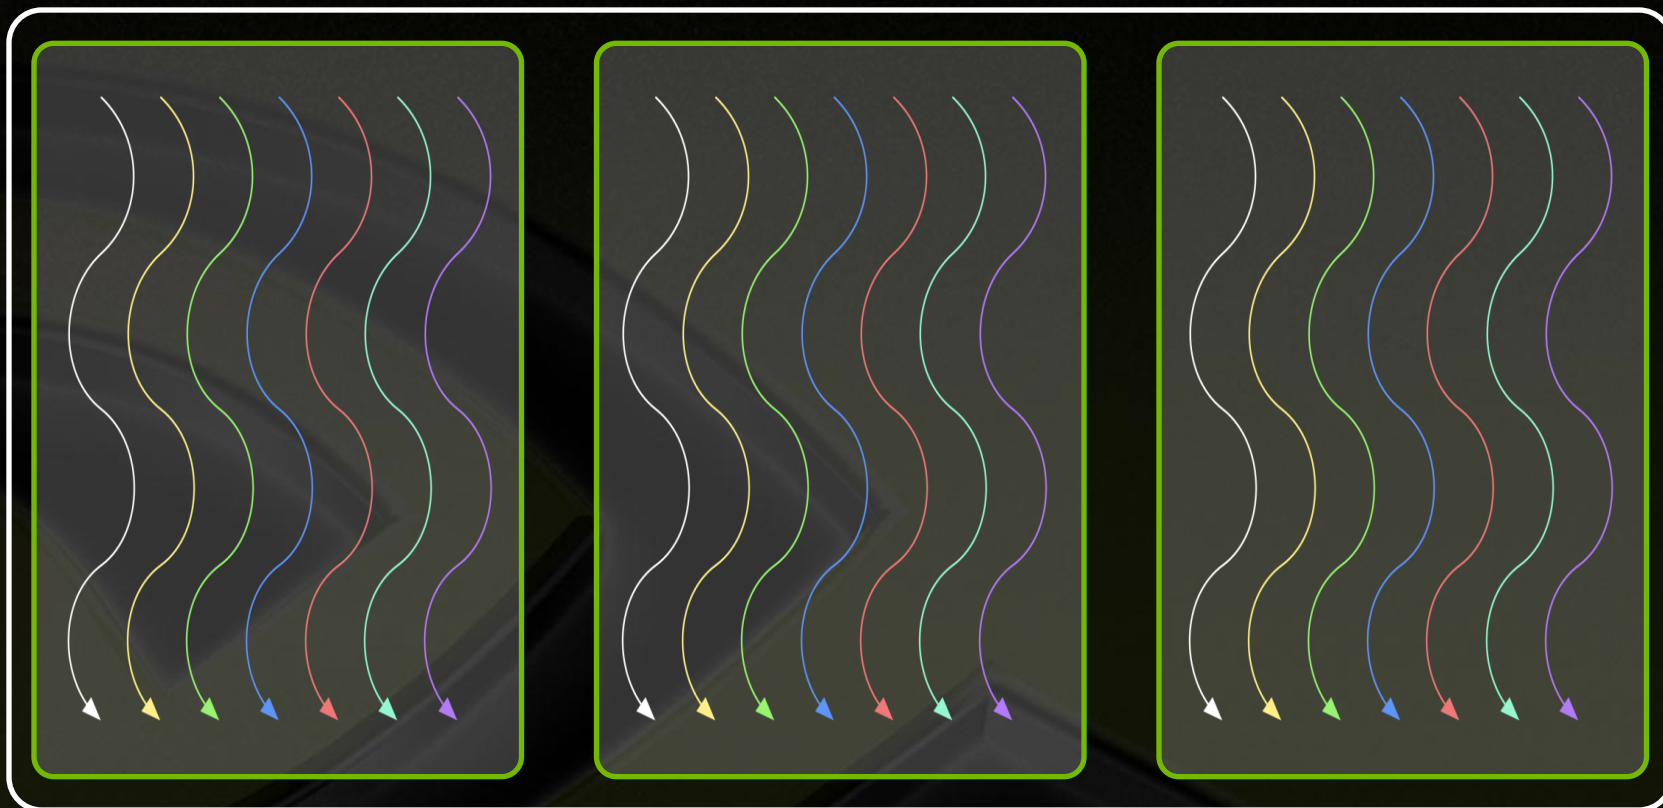


CUDA Kernels: Subdivide into Blocks



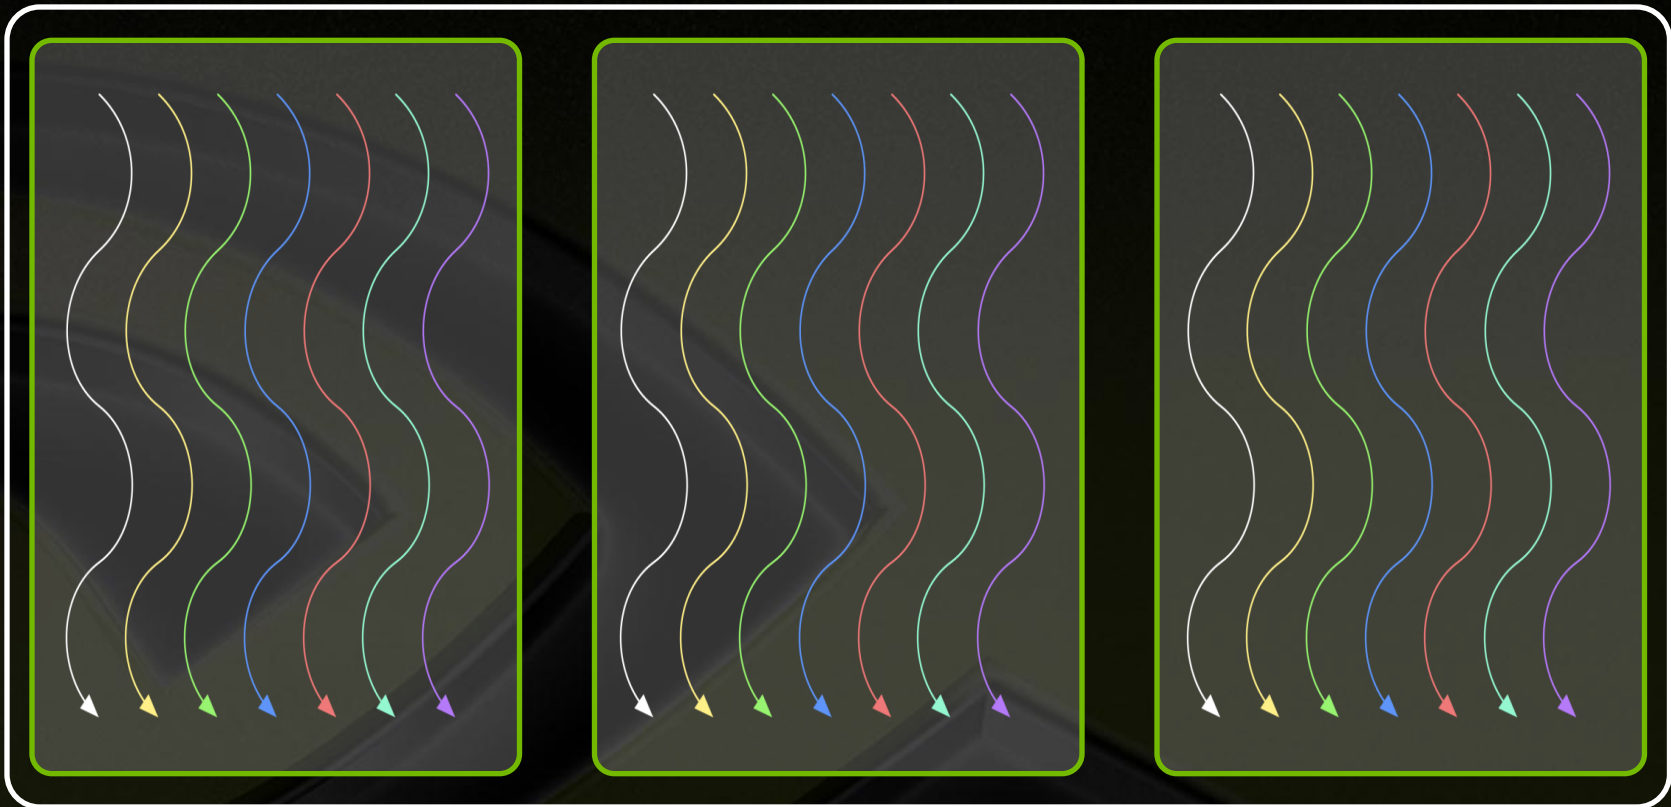
- Threads are grouped into **blocks**

CUDA Kernels: Subdivide into Blocks




- Threads are grouped into **blocks**
- **Blocks** are grouped into **a grid**

CUDA Kernels: Subdivide into Blocks



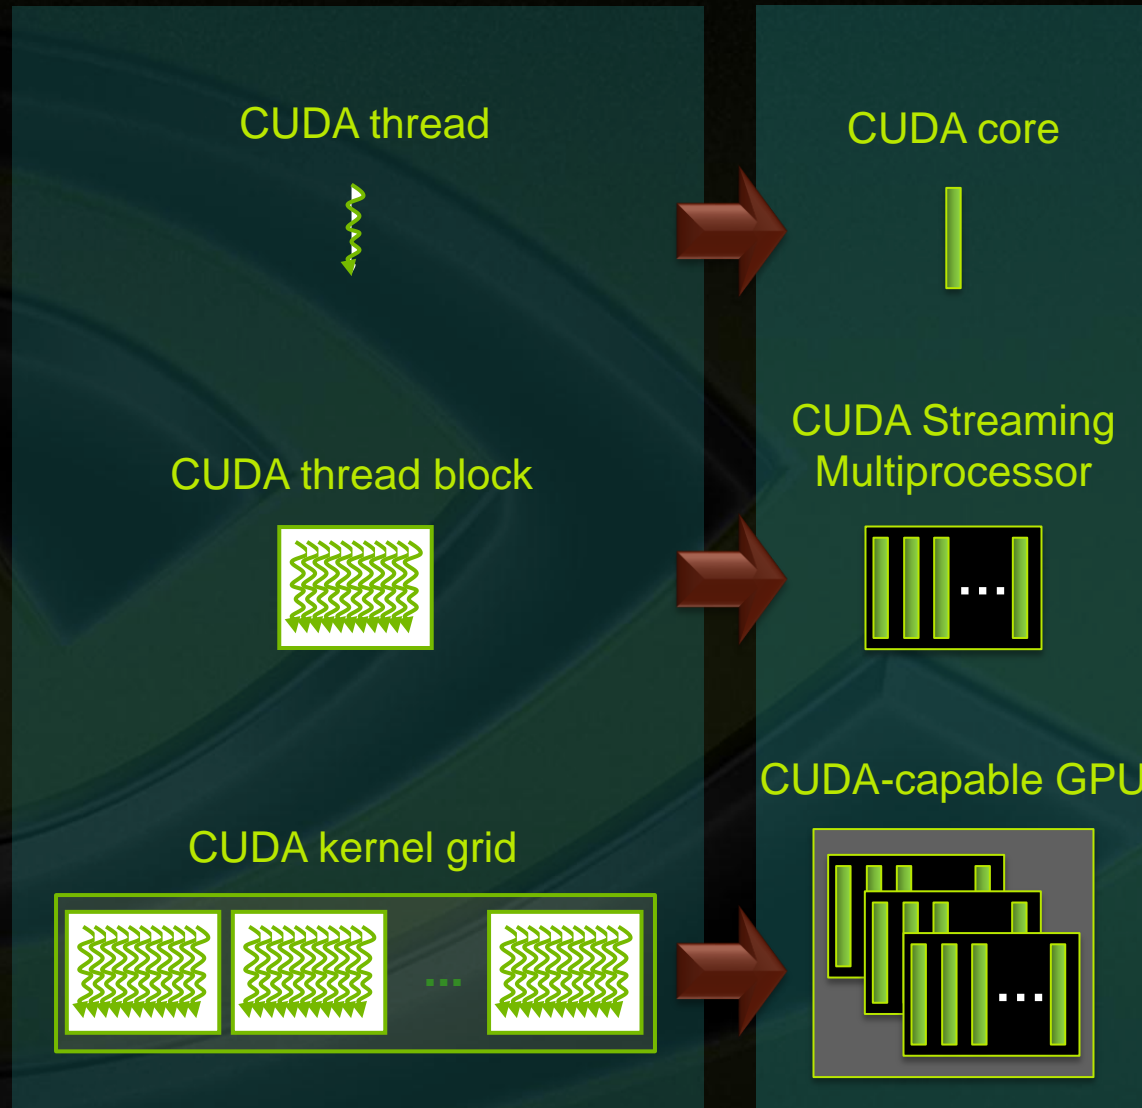
- Threads are grouped into **blocks**
- **Blocks** are grouped into **a grid**
- A **kernel** is executed as a **grid of blocks of threads**

CUDA Kernels: Subdivide into Blocks



- Threads are grouped into **blocks**
 - Note: Adjacent threads execute in lock-step scheduling groupings called **warps**; a block comprises one or more warps
- **Blocks** are grouped into a **grid**
- A **kernel** is executed as a **grid** of **blocks** of **threads**

Kernel Execution



- Each thread is executed by a core
- Each block is executed by one SM and does not migrate
- Several concurrent blocks can reside on one SM depending on the blocks' memory requirements and the SM's memory resources
- Each kernel is executed on one device
- Multiple kernels can execute on a device at one time

Thread blocks allow cooperation

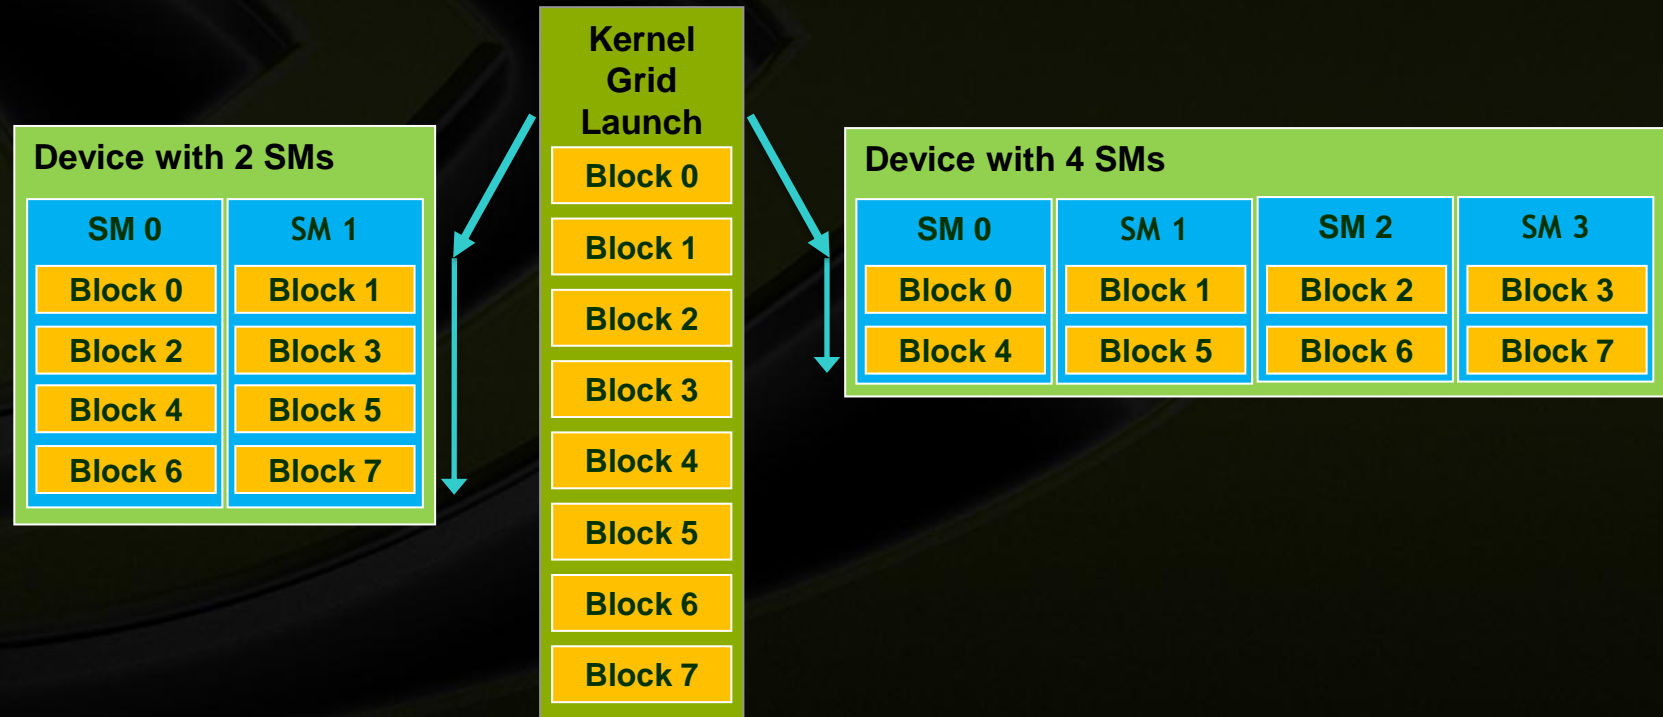


- **Threads may need to cooperate:**
 - **Cooperatively load/store blocks of memory that they all use**
 - **Share results with each other or cooperate to produce a single result**
 - **Synchronize with each other**

Thread blocks allow scalability



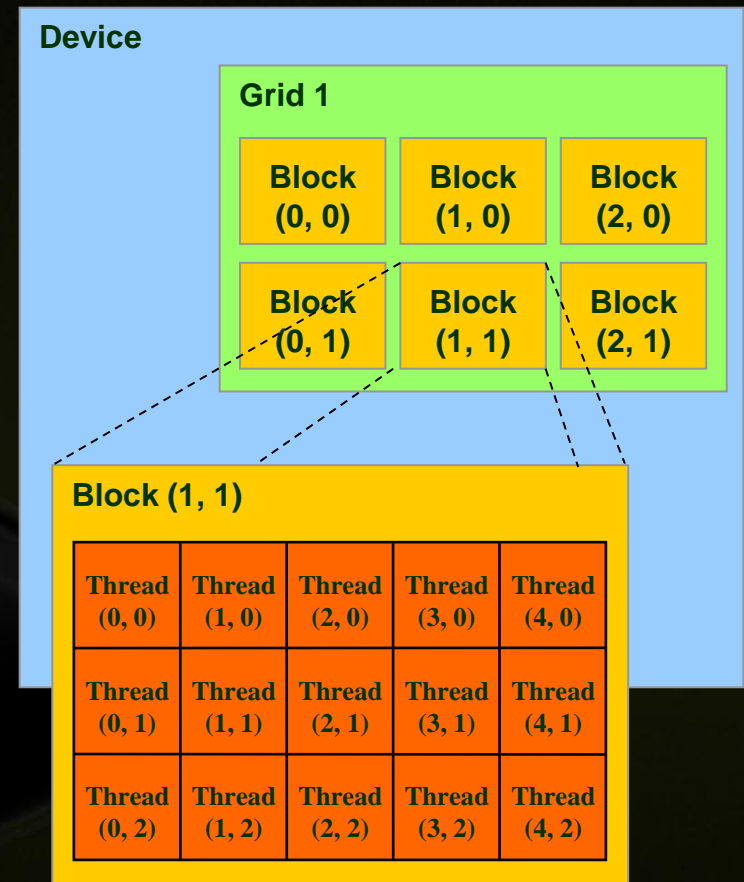
- Blocks can execute in any order, concurrently or sequentially
- This independence between blocks gives scalability:
 - A kernel scales across any number of SMs



IDs and Dimensions



- **Threads:**
 - 3D IDs, unique within a block
- **Blocks:**
 - 2D IDs, unique within a grid
- **Dimensions set at launch time**
 - Can be unique for each section
- **Built-in variables:**
 - threadIdx, blockIdx
 - blockDim, gridDim



Launching kernels

- Modified C function call syntax:

```
kernel<<<dim3 grid, dim3 block>>>(...)
```

- Execution Configuration (“<<< >>>”):

- grid dimensions: **x** and **y**
- thread-block dimensions: **x**, **y**, and **z**

```
dim3 grid(16, 16);  
dim3 block(16,16);  
kernel<<<grid, block>>>(...);  
kernel<<<32, 512>>>(...);
```

Minimal Kernels

```
__global__ void minimal( int* d_a)
```

```
{
```

```
    *d_a = 13;
```

```
}
```

```
__global__ void assign( int* d_a, int value)
```

```
{
```

```
    int idx = blockDim.x * blockIdx.x + threadIdx.x;
```

```
    d_a[idx] = value;
```

```
}
```

Common Pattern!

Example: Increment Array Elements



Increment N-element vector a by scalar b



Let's assume $N=16$, $\text{blockDim}=4 \rightarrow 4$ blocks



$\text{blockIdx.x}=0$
 $\text{blockDim.x}=4$
 $\text{threadIdx.x}=0,1,2,3$
 $\text{idx}=0,1,2,3$

$\text{blockIdx.x}=1$
 $\text{blockDim.x}=4$
 $\text{threadIdx.x}=0,1,2,3$
 $\text{idx}=4,5,6,7$

$\text{blockIdx.x}=2$
 $\text{blockDim.x}=4$
 $\text{threadIdx.x}=0,1,2,3$
 $\text{idx}=8,9,10,11$

$\text{blockIdx.x}=3$
 $\text{blockDim.x}=4$
 $\text{threadIdx.x}=0,1,2,3$
 $\text{idx}=12,13,14,15$

$\text{int idx} = \text{blockDim.x} * \text{blockIdx.x} + \text{threadIdx.x};$
will map from local index threadIdx to global index

NB: blockDim should be ≥ 32 in real code, this is just an example

Example: Increment Array Elements



CPU program

```
void increment_cpu(float *a, float b, int N)
{
    for (int idx = 0; idx < N; idx++)
        a[idx] = a[idx] + b;
}
```

```
void main()
{
    .....
    increment_cpu(a, b, N);
}
```

CUDA program

```
__global__ void increment_gpu(float *a, float b, int N)
{
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx < N)
        a[idx] = a[idx] + b;
}
```

```
void main()
{
    .....
    dim3 dimBlock (blocksize);
    dim3 dimGrid( ceil( N / (float)blocksize) );
    increment_gpu<<dimGrid, dimBlock>>>(a, b, N);
}
```

Minimal Kernel for 2D data

```
__global__ void assign2D(int* d_a, int w, int h, int value)
{
    int iy = blockDim.y * blockIdx.y + threadIdx.y;
    int ix = blockDim.x * blockIdx.x + threadIdx.x;
    int idx = iy * w + ix;

    d_a[idx] = value;
}

...
assign2D<<<dim3(64, 64), dim3(16, 16)>>>(...);
```

CUDA Memory Model



Memory Model

- **Registers**

- Per thread
- Data lifetime = thread lifetime

- **Local memory**

- Per thread off-chip memory (physically in device DRAM)
- Data lifetime = thread lifetime

- **Shared memory**

- Per thread block on-chip memory
- Data lifetime = block lifetime


- **Global (device) memory**

- Accessible by all threads as well as host (CPU)
- Data lifetime = from allocation to deallocation


- **Host (CPU) memory**

- Not directly accessible by CUDA threads

Memory Model



Memory Model



Memory Model



Introduction to CUDA Programming



Outline of CUDA Basics



- **Basics to setup and execute CUDA code:**
 - GPU memory management
 - Extensions to C++ for kernel code
 - GPU kernel launches
- **Some additional basic features:**
 - Checking CUDA errors
 - CUDA event API
- **See the Programming Guide for the full API**
<http://docs.nvidia.com/cuda/cuda-c-programming-guide>

GPU Memory Allocation / Release



- **Host (CPU) manages GPU memory:**
 - **cudaMalloc (void ** pointer, size_t nbytes)**
 - **cudaMemset (void * pointer, int value, size_t count)**
 - **cudaFree (void* pointer)**

```
int n = 1024;  
int nbytes = 1024*sizeof(int);  
int * d_a = 0;  
cudaMalloc( (void**)&d_a, nbytes );  
cudaMemset( d_a, 0, nbytes);  
cudaFree(d_a);
```

Data Copies



- **cudaMemcpy(void *dst, void *src, size_t nbytes, enum cudaMemcpyKind direction);**
 - returns after the copy is complete
 - blocks CPU thread
 - doesn't start copying until previous CUDA calls complete
- **enum cudaMemcpyKind**
 - cudaMemcpyHostToDevice
 - cudaMemcpyDeviceToHost
 - cudaMemcpyDeviceToDevice
- **Non-blocking memcpyes are provided**

Programming Exercise 1



- Send a string to the GPU and back
- Allocate GPU memory for n chars
- Copy from CPU to GPU and back
- Print the values

Programming Exercise 1



- **Compiling a CUDA application**

```
nvcc --arch=sm_35 -o sample1 sample1.cu
```

- **Execution**

```
salloc --res=course --nodes=1 --walltime=00:10:00
```

```
aprun -n 1 ./sample1
```

- **Some CUDA API**

```
cudaMalloc (void ** pointer, size_t nbytes)
```

```
cudaFree (void* pointer)
```

```
cudaMemcpy( void *dst, void *src, size_t nbytes,  
            enum cudaMemcpyKind direction);
```

Programming Exercise 1



```
#include <stdio.h>
#define N 100

int main(int argc, char** argv){
    char msg_in[] = "hello accelerator";
    char msg_out[N];
    char* d_msg;                // will be message on the device

    cudaMalloc(&d_msg, N*sizeof(char));
    cudaMemcpy(d_msg, msg_in, N*sizeof(char), cudaMemcpyHostToDevice);

    cudaMemcpy(msg_out, d_msg, N*sizeof(char), cudaMemcpyDeviceToHost);

    printf("msg_out= %s\n", msg_out);
    cudaFree(d_msg);

    return 0;
}
```

Code executed on GPU



- **C++ function with some restrictions:**
 - Can only access GPU memory (with some exceptions)
 - No variable number of arguments
 - No static variables
- **Must be declared with a qualifier:**
 - **__global__** : launched by CPU,
cannot be called from GPU
must return void
 - **__device__** : called from other GPU functions,
cannot be launched by the CPU
 - **__host__** : can be executed by CPU
 - **__host__** and **__device__** qualifiers can be combined
- **Built-in variables:**
 - **gridDim, blockDim, blockIdx, threadIdx**

Variable Qualifiers (GPU code)

- **__device__**
 - stored in global memory (not cached, high latency)
 - accessible by all threads
 - lifetime: application
- **__constant__**
 - stored in global memory (cached)
 - read-only for threads, written by host
 - Lifetime: application
- **__shared__**
 - stored in shared memory (latency comparable to registers)
 - accessible by all threads in the same threadblock
 - lifetime: block lifetime
- **Unqualified variables:**
 - Stored in local memory:
 - scalars and built-in vector types are stored in registers
 - arrays are stored in device memory

Launching kernels on GPU

- **Launch parameters:**
 - grid dimensions (up to 2D)
 - thread-block dimensions (up to 3D)
 - shared memory: number of bytes per block
 - for extern smem variables declared without size
 - Optional, 0 by default
 - stream ID
 - Optional, 0 by default

```
dim3 grid(16, 16);
```

```
dim3 block(16,16);
```

```
kernel<<<grid, block, 0, 0>>>(...);
```

```
kernel<<<32, 512>>>(...);
```

Programming Exercise 2



- **Build on Exercise 1**
- **Write a kernel to convert a simple string to upper case (Fact: 'a' = 96, 'A' = 64)**

Programming Exercise 2



```
#include <string.h>
#include <stdio.h>
#define N 100

__global__ void kernel(char* msg, int n){
    int tid = threadIdx.x;
    if(threadIdx.x < n) {
        char m = msg[tid];
        msg[tid] = m > 96 ? m - 32 : m ;
    }
}

int main(int argc, char** argv){
    ...
    cudaMemcpy(d_msg, msg_in, N*sizeof(char), cudaMemcpyHostToDevice);
    kernel<<<1, 100>>>>(d_msg, strlen(msg_in));
    cudaMemcpy(msg_out, d_msg, N*sizeof(char), cudaMemcpyDeviceToHost);
    ..
}
```

Thread Cooperation via Smem

- Shared memory accessible by all threads in block

```
__shared__ double a[32];
```

- “software managed cache”
- Shared memory partitioned among all blocks running on a SM

Thread Synchronization Function



- `void __syncthreads ();`
- **Synchronizes all threads in a block**
 - Once all threads have reached this point, execution resumes normally
 - Used to avoid RAW / WAR / WAW hazards when accessing shared memory
- **Should be used in conditional code only if the conditional is uniform across the entire thread block**

Programming Exercise 3



- **Build on Exercise 2**
- **Write a kernel to convert the first letter in a word to upper case**

Programming Exercise 3



```
__global__ void kernel(char* msg, int n){

    __shared__ char s_msg[N];
    int tid = threadIdx.x;

    if(tid < n)    s_msg[tid] = msg[tid];

    __syncthreads();

    if(tid < n & tid > 0) {
        char m = s_msg[tid - 1];
        if(m == ' ') msg[tid] = s_msg[tid] - 32;
    }
}
```

Host Synchronization



- **All kernel launches are asynchronous**
 - control returns to CPU immediately
 - kernel starts executing once all previous CUDA calls have completed
- **Memcopies are synchronous**
 - control returns to CPU once the copy is complete
 - copy starts once all previous CUDA calls have completed
- **`cudaThreadSynchronize()`**
 - blocks until all previous CUDA calls complete
- **Asynchronous CUDA calls provide:**
 - non-blocking memcopies
 - ability to overlap memcopies and kernel execution

Device Management



- **CPU can query and select GPU devices**
 - `cudaGetDeviceCount(int* count)`
 - `cudaSetDevice(int device)`
 - `cudaGetDevice(int *current_device)`
 - `cudaGetDeviceProperties(cudaDeviceProp* prop, int device)`
 - `cudaChooseDevice(int *device, cudaDeviceProp* prop)`
- **Multi-GPU setup:**
 - device 0 is used by default

CUDA Error Reporting to CPU



- **All CUDA calls return error code:**
 - except for kernel launches
 - `cudaError_t` type
- **`cudaError_t cudaGetLastError(void)`**
 - returns the code for the last error (no error has a code)
- **`char* cudaGetErrorString(cudaError_t code)`**
 - returns a null-terminated character string describing the error

```
printf("%s\n", cudaGetErrorString( cudaGetLastError() ) );
```

CUDA Event API



- Events are inserted (recorded) into CUDA call streams
- Usage scenarios:
 - measure elapsed time for CUDA calls (clock cycle precision)
 - query the status of an asynchronous CUDA call
 - block CPU until CUDA calls prior to the event are completed
 - **asyncAPI** sample in CUDA SDK

```
cudaEvent_t start, stop;  
cudaEventCreate(&start);          cudaEventCreate(&stop);  
cudaEventRecord(start, 0);  
kernel<<<grid, block>>>(...);  
cudaEventRecord(stop, 0);  
cudaEventSynchronize(stop);  
float et;  
cudaEventElapsedTime(&et, start, stop);  
cudaEventDestroy(start); cudaEventDestroy(stop);
```

What about Fortran?



- **Mixed language approach**

```
call myfun()  
extern "C" void myfunc() {  
    mykernel<<< grid, block >>>()  
}
```

- **PGI's CUDA Fortran**

```
real*8, dimension(:) :: u  
real*8, device, dimension(:) :: a, b  
a = u  
call mykernel<<< grid, block >>> ( a, b)  
  
attributes(global) subroutine mykernel(a, b)  
..  
end subroutine mykernel
```


What did we skip? What's next?



- **Concurrent kernels and data transfers**
 - Streams
- **Optimization**
 - Global memory access, “coalescence”
- **Mapped memory**
- **Multi-GPU programming**
 - Device-device transfers